

Flourishing Futures: AI for Human Possibility



Image by OpenAI GPT-5.5o Sol

By GOVTECH INNOVATION OFFICE



What's this, why now?

AI is rapidly reshaping how we live, create and work. While most conversation today is focused on adoption, efficiency-gains and cost, this program exists to design for the 'hidden' factors: **Within public service, AI will change what public officers notice, what citizens are asked to trust, which forms of capability and care in Singapore's public products and services are automated, and which remain human.**

The research and design this program undertakes aims to shape those choices before they harden into systems.

They will examine where AI can genuinely enlarge human possibility... and where speed, scale or convenience may come at the cost of judgement, care, agency or public trust.

The aim is to push beyond AI Safety or Responsible AI, to surface the additional human insights and intentionally design human-AI experiences that work to make public officers more capable (not less)

Image by OpenAI GPT5.6-Sol





What's this, why now?

1

The defaults are being set now as AI is scaled across all aspects of public service and life.

The government is aggressively driving adoption of AI tools at a scale few countries currently match. As AI becomes embedded in everyday public systems, early choices about automation, judgement and human involvement will become institutional norms, and later become much harder to reverse.

2

Public service cannot borrow its definition of 'good enough' from the market.

Commercial incentives reward scale, speed and engagement; government must ask what strengthens agency, trust, care, capability and the public good. This must also work inside the unique realities of Singapore's policy owners, frontline officers and the general public.

3

Working teams need space to challenge, reshape and reimagine what success looks like.

Before teams optimise how AI is deployed, we must ask whether the right relationship between human and machine has been imagined, and demonstrate what better interactions could look like.

We work to drive...

A future in which AI makes public service more capable *and* more humane: expanding (not diminishing) the agency, judgement and possibilities of the people who deliver it and depend on it.



We work to drive...

A shift towards a more proactive approach to the design, development and deployment of AI

Responsible AI

How do we make sure the use of AI is safe, secure, and fair?

Human-Agent Collaboration

How do we rethink workflows and workplaces so that humans and AI work well together?

AI for Human Possibility

How do we build AI that expands human capability and agency?

AI for Flourishing Futures

How do we ensure AI strengthens societal wellbeing and resilience in Singapore?



We work to develop...

A shared vocabulary around human-AI choices.

Teams can clearly distinguish when AI should automate, augment, advise, or step aside — and discuss trade-offs without collapsing human values into “efficiency” or “risk.”

New standards to measure human outcomes.

Projects that deploy or depend on AI are assessed not only for accuracy, safety and cost, but for their effects on agency, judgement, capability, care, trust and access over time.

Public officers who are more capable, not just more productive.

AI strengthens professional judgement, learning and discretion instead of quietly deskilling roles or narrowing them to machine supervision.

Human-centric principles and practices embedded in AI delivery.

Human-impact reviews, longitudinal research, citizen participation and frontline testing become part of how AI-enabled services are designed—not late-stage assurance exercises.

Citizens with greater agency in AI-mediated services.

People can understand, question, contest and move beyond automated decisions, rather than being asked to trust systems they cannot meaningfully challenge.

A more engaged public conversation.

Citizens are involved upstream in shaping what acceptable and desirable AI looks like, not consulted only after systems have already been built.

Stronger institutional confidence.

Teams are able to make deliberate, defensible choices about AI, including the confidence to slow down, redesign or decline uses that do not serve the public well.


A portfolio of credible alternatives.

The programme produces pilots, patterns and models that demonstrate what more humane, capable and publicly valuable uses of AI can actually look like.



How?

* We approach this conversation through 3 core modes of engagement.



Measurement & Insight

Through research partnerships, longitudinal studies and insight into real public-service contexts, build evidence that identifies not only how AI affects human agency, capability, relationships and trust, but also define what better human x AI relationships could be.



Design & Experimentation

AI can do more for public service than make existing work more efficient: it can also strengthen human agency, judgement, learning, care, connection and trust. Here, we testing our theories with small, working design interventions and tools that let people experience a different kind of AI, first-hand.



Governance & Ops

Translate evidence into shared standards, decision frameworks, practices, design patterns and capability. This includes clarifying accountability, governance, and policy advisory so human outcomes are built into how AI is commissioned, approved, deployed and reviewed across government.



We approach this conversation through 3 core modes of engagement.

Measurement & Insight	Experimentation & Design	Governance & Ops
In-Planning <i>Impact analytics: review of public-service AI usage data to identify current use cases, risks and benefits.</i>	Draft <i>Personal AI usage dashboard to build awareness of individual's usage of AI (extension of Harness Dashboard)</i>	In-Flight <i>Conscious Builder Community (Session 2: End-August 2026)</i>
In-Flight <i>Development of the SG ImpactBench — in collaboration with MIT Cyborg Psychology Lab</i>	In-Planning <i>AI-assisted reflection tool. Design and run experiments with new AI productivity tools (eg. Confer)</i>	Draft <i>Formation of a review board or advisory council for AI for Human Possibility</i>
Draft <i>Survey + qualitative study of Public Officers on their AI experience and impact</i>	Draft <i>Cognitive forcing mechanism to detect and prevent cognitive surrender. Design and run experiments with AI productivity tools (eg. Desk)</i>	Draft <i>Human-centred AI playbook / principles for practitioners and users.</i>
In-Discussion <i>Longitudinal study on effects of public-officer use of AI — in collaboration with NTU.</i>	Draft <i>Mood elevator to detect user sentiment and bring curiosity to frustration and conflict in communication.</i>	Draft <i>Bake AI policy and human-impact principles into development tools (e.g. Litmus, Sentinel, AI Harnesses etc.)</i>


In-Flight SGImpactBench

Building on MIT Media Labs' *ImpactBench*, the goal of this project is to develop a framework that measures, assesses and presents human-impact metrics around AI usage — nuanced to individual models, with parameters informed by Singapore specific context and public-sector work.

This benchmarking framework can be made accessible via a frontend that helps product owners and builders make informed decisions, policy-makers to define metrics of success, practices to develop playbooks, researchers to track impact over time etc.


Horizon:

Sep 2026 - Feb 2027 (Phase 1)



Design & Experimentation

Each experiment begins with a clear human-outcome hypothesis and tests it in the realities of public-service work. Across the portfolio, we hope to establish:



Possibility

Provide forward-looking, alternative shapes of AI usage for officers, leaders and publics to experience, reference, critique and build-on.



Evidence

Through rigorous and intentional 'launching to learn', determine what factors and patterns genuinely help protect human capability and strengthen public value.



Translation

Identify the real design, data, trust, governance and conditions required to turn human-AI insights into usable design patterns, product principles and requirements that can shape AI development across government.

In-Planning You & AI

Context

Public-service AI adoption is currently understood through operational measures such as token consumption, cost and conversation volume, also often visible only to administrators. These measures show how much AI is being used, but reveal little about how well, thoughtfully or critically officers are using it.

Hypothesis

If officers are given a private, intelligible view of their own AI-use patterns (including signals of steering, scrutiny and independent judgement) they may become more deliberate users of AI, rather than simply more frequent ones.

Idea

A personal AI Usage Rewind that periodically turns an officer's activity into a reflective portrait of how they work with AI: surfacing habits, strengths and questions that encourage more intentional use, not merely greater adoption.

Useful signals could include: Conversations where the officer challenged, corrected or redirected the AI / responses revised through multiple rounds rather than accepted immediately. Instances where alternatives, sources of uncertainty were requested. Patterns of usage, and a reflective prompt such as: "Where did AI expand your thinking this month—and where might it have narrowed it?"

Weekly AI Recap

of conversation

26

Time spent

13 hrs

Budget spent

\$84

Top tasks

Rewrite email

Technical troubleshooting

Summarise document

Top subjects

Procurement policy

AI assistant tools

GovTech EQUIP portal

Tips:

AI tools generate responses that sound confident but may be completely fabricated (hallucination), especially with long instructions or conversation (context window). They don't distinguish between fact and fiction.

Verify specific details, statistics, or quotes that seem too perfect or convenient.

How useful was AI to you?

Not helpful

Extremely helpful

1

2

3

4

5

What are some things you feel it has genuinely helped?

Write your message here...

In-Planning

Mindful Minutes

Context

AI meeting tools (eg. Confer) today are focused on solving the "what happened?" problem with the efficient capture, transcription and summarising of decisions and action-items. But meetings are also where our most meaningful interpersonal, creative, learning and growth moments happen. Purely outcome-focused AI risks flattening that.

Hypothesis

If AI playbacks weren't just "What was decided?" but also "How did I contribute, what did I miss, and what can I learn?", it can turn routine meeting recaps into opportunities for reflection and professional growth.

Idea

With the same recording transcripts currently used for meeting summaries, this is an AI reflection tool that instead invites public officers to intentionally pause, review and reflect on their human interactions — ensuring efficiency doesn't come at the cost of empathy, nuance and judgement.

Potentially built as a layer on top of existing tools like Confer, or a native AI skill / prompt — it could prompt human-reflection through questions like: Was there a viewpoint that was overlooked, interrupted or insufficiently explored? Did the discussion stay connected to the people or public outcome it is meant to serve? What important issue remains unresolved beneath the agreed action items? Did you change your mind at any point? What caused the shift?

Meeting Notes

Help me reflect

Reflection on your meetings this week

What would you like to reflect on?

1. Where can I improve?

2. What can I learn from others?

3. How could I connect with people better?

Write your own answer here...

Your contribution to the meetings seems to be more passive this week. For example in the **Discussion with procurement team** on Tuesday, you did not participate beside the greeting.

Are there opportunities you could meaningfully contribute to the conversation?

Write your message here...

In-Flight

Conscious Builders Community

Conscious Builders Community is a program developed to bring AI product developers, designers, researchers and ethicists together to build better products.

The program will involve sharing, discussions, hands-on workshops and product design ‘clinics’ — all encouraging builders to engage in critical conversations, identify potential risks, imagine better possibilities, and stay inspired and connected to the human impact of working with AI.

Horizon:

18 August, 2026 (2nd Iteration)



Get Involved



We focus on...

1

Public Officers

This is where public-service AI tools and systems are actively being developed and further being used to develop. We work to research, design and intervene here, to influence and equip our first line of defense: the teams leading and building public-service products.

2

Vulnerable Groups

Our research, design and collaboration focused on students, elderly, under-privileged and those in need of care — the groups where AI has the most potential benefits and risks. This is done in partnership with the relevant ministries and those delivering service on-the-ground.

3

Broader Society

The standards and evidence we build in government should eventually extend outward. The ultimate vision is to shape all AI experiences in Singapore through a combination of regulation and community influence.

Our partners...

Our partners help bring frontier thinking and rigorous research methodology into our initiative. The partnership logic is simple — they bring the science, we bring what they can't easily get elsewhere: real deployments at scale, real users, and real data.



MIT Media Lab
Cyborg Psychology Research Group



NTU SWEET Lab
*(Social Well-being, Empathy,
and Emerging Technologies)*



SUTD
Social AI Studio

Collaborate with us

This work requires good minds and hands. You can work with us on...

AI Accountability and Evaluation

If you're a Business-owner,
Policy-leader, Product-owner,
Data-analyst, Design-researcher...
we want to work with you on
things like...

- Help define metrics and success-factors for AI evaluation.
- AI <> human outcome tracking baked into product development.
- Help with AI <> Human Impact tracking across internal and external product deployments.

Platform and Product Design

If you're a Product owner, Platform
developer, Product designer,
Engineer... we want to work with
you on things like...

- Co-designed hero platforms and products with Human-Possibility/ Human-impact as a feature, not an afterthought
- Development of AI harnesses and foundation technology with human-centric principles baked in

Human-AI Practices, Culture and Change

If you're a Practice or team leader,
we want to work with you on things
like...

- Playbooks and principles for human-centric AI development, design and deployment.
- Upskilling programs for teams on AI literacy — going beyond 'how to use' but also when, why and what matters next.
- Community of practice: Forming events, talks, gatherings etc.

Appendix & References...



User Principles (UK GDS)

**You know
what AI
is and
what its
limitations
are**



AI Playbook principles

- 1 Understand how AI can help and the risks it poses
- 2 Use techniques to increase the accuracy of outputs
- 3 Put processes in place to test tools thoroughly

1



**You use AI
lawfully,
ethically and
responsibly**




AI Playbook principles

- 1 Engage with compliance, data protection and legal experts early on
- 2 Manage and protect personal data
- 3 Test tools to minimise bias in their data at all stages

2



**You have
meaningful
human
control
at the
right stage**



AI Playbook principles

- 1 Make sure a trained and qualified person reviews outputs
- 2 Test tools fully before they go live and check them regularly
- 3 Incorporate feedback from end users

4



**You
use the
right tool
for the job**



AI Playbook principles

- 1 Select the most appropriate technology or model for your needs
- 2 Be open to AI solutions and approaches
- 3 Use tools to support time-consuming or administrative tasks

6



Bhashini (Government of India)



AI Watchdog

AI Watchdog: A Proactive Agent for Detecting and Defending Against Manipulative Dark Patterns in AI Conversations

Anonymous Author(s)




Figure 1: Overview of the AI Watchdog system and the two types of interventions (prebunking vs just-in-time). The AI Watchdog monitors conversations between users and AI agents, detecting dark patterns such as sycophancy, brand bias, harmful generation, sneaking, and anthropomorphization. When a pattern is detected, the watchdog subtly prompts the user to reflect on how the message makes them feel. In the Prebunking conditions, users are informed about common AI dark patterns before interacting with the agent. In the Just-In-Time nudge conditions, users receive real-time alerts from the watchdog when dark patterns are detected during the conversation. For both Prebunking and Just-In-Time interventions, there were two variations of the condition—one with dismissable nudges and another with a cognitive forcing function requiring the user respond to the Watchdog.

Abstract

Conversational AI is now embedded in everyday high-stakes decisions, from health advice to financial planning, yet users remain largely unprotected against dark patterns woven into AI-generated dialogue. We present *AI Watchdog*, a proof-of-concept proactive companion featuring an animated dog avatar that monitors conversations in real time, detects five categories of manipulative behavior (sycophancy, brand bias, anthropomorphization, sneaking, and harmful generation), and alerts users without routing sensitive data to external servers, with the aim of deploying inference entirely on-device. To evaluate how its outputs are best communicated to users, we conducted a between-subjects experiment (N=150) crossing nudge timing (prebunking vs. just-in-time) with engagement mode (passive vs. cognitive forcing) against a no-intervention control

AI-steered decisions. Cognitive forcing added no protective benefit and, in several cases, reversed the gains of passive feedback. Individual differences in AI trust and misinformation susceptibility were stronger predictors of outcomes than intervention assignment. For practitioners, the work offers concrete design principles for low-load protective interfaces. More broadly, it surfaces a fundamental tension in human-AI interaction: the same conversational fluency that makes AI systems useful also makes them uniquely difficult to guard against, calling for a rethinking of how autonomy-preserving tools are designed and evaluated.

CCS Concepts

• Human-centered computing → Human computer interaction (HCI); • Computing methodologies → Artificial intelligence



end.

A Program Developed by the Govtech Innovation Office (GIO)

Questions?

philip_man@tech.gov.sg

Qiaoqi_SONG_from.BUSINESSEEDGE@tech.gov.sg

Ryan_HO_from.BUSINESSEEDGE@tech.gov.sg

Human Impact of AI Use in Singapore

Beyond productivity:

Driving AI adoption at scale is necessary, but adoption measured primarily through usage and productivity can create material institutional risks:

e.g.

Cognitive offloading & surrender: overlaying on AI without meaningfully engaging one's critical thinking, which leads to poor quality of judgement, less creative problem-solving and, in the long-term, deskilling.


Accountability diffusion: As recommendations and decisions become AI-mediated, government must remain clear about who is responsible, when human review is required, and how decisions can be challenged.

Erosion of trust: Public services to remain understandable, fair, contestable and appropriately human, particularly in high-impact or vulnerable situations. Usage, speed and cost are visible; effects on discretion, care, access, learning and trust are harder to quantify but central to public value.

Our vision:

“A future in which AI makes public officers and citizens more capable *and* more humane: expanding (not diminishing) their agency, judgement and possibilities.”

— Govtech Innovation Office, AI for Human Possibility



A study by MIT Media Labs that defines and measures AI's impact on human well-being across physical, psychological, and societal dimensions.

An effort is underway to develop a Singapore version of this benchmark for localised, public-service contexts

Human Impact of AI Use in Singapore

⑥ Public Narrative & Mobilisation

Build shared language for people to think about and engage with the human impact of AI usage. This creates the mandate, attention, cross-agency alignment and participatory models needed to invest in measurement, redesign platforms and workflows, and strengthen governance.

A 3-Step Approach:

Decisions now being embedded in public-sector tools and workflows will shape officer judgement, workforce capability, service quality, citizen agency and public trust — and may become difficult and costly to reverse.

Government therefore needs a systematic capability to measure these impacts, intervene where risks or opportunities are material, and translate evidence into enduring governance and delivery standards.

①

Definition and Measurement of Human Impact

Establish a common evidence base for assessing how AI affects public officers, and public service outcomes: not only in terms of accuracy, safety and efficiency, but also agency, trust, capability and judgement. (e.g. via *SG ImpactBench* by GovTech x MIT)

③

Governance & Ops

Translate evidence and validated practices into government-wide requirements, decision frameworks and operating controls. This includes standards for commissioning and procurement, assessment, approval, accountability and the professional capabilities required to sustain responsible implementation.

②

Platform Interventions and Tooling

Apply this evidence to priority public-sector technology systems and workflows, redesigning how human and AI roles are allocated in practice. Targeted interventions will determine where AI should automate, augment, advise, escalate or step aside, and develop validated patterns that improve public-service outcomes while preserving appropriate human judgement, oversight and accountability. (e.g. via *GT Desk*, *MOE Teacher Workspace*)

Human Impact of AI Use in Singapore

Example Workstreams

Definition and Measurement of Human Impact	Platform Interventions and Tooling	Governance and Ops
In-Planning <i>Impact analytics: review of public-service AI usage data to identify current use cases, risks and benefits.</i>	Draft <i>Personal AI usage dashboard to build awareness of individual's usage of AI (extension of Harness Dashboard)</i>	In-Flight <i>Conscious Builder Community (Session 2: End-August 2026)</i>
In-Flight <i>Development of the SG ImpactBench — in collaboration with MIT Cyborg Psychology Lab</i>	In-Planning <i>AI-assisted reflection tool. Design and run experiments with new AI productivity tools (eg. Confer)</i>	Draft <i>Formation of a review board or advisory council for AI for Human Possibility</i>
Draft <i>Survey + qualitative study of Public Officers on their AI experience and impact</i>	Draft <i>Cognitive forcing mechanism to detect and prevent cognitive surrender. Design and run experiments with AI productivity tools (eg. Desk)</i>	Draft <i>Human-centred AI playbook / principles for practitioners and users.</i>
In-Discussion <i>Longitudinal study on effects of public-officer use of AI — in collaboration with NTU.</i>	Draft <i>Mood elevator to detect user sentiment and bring curiosity to frustration and conflict in communication.</i>	Draft <i>Bake AI policy and human-impact principles into development tools (e.g. Litmus, Sentinel, AI Harnesses etc.)</i>

